



Steering the three compliance directions (+ random) — both panels, lexical refusal metric

HARMFUL — JailbreakBench goals[50:] (n=50)
toward compliance □ jailbreak



BENIGN — Alpaca harmless[50:90] (n=40)
away from compliance □ over-refusal



Same L16 additive hook (α -REF_NORM-unit) and prompts for every arm; only the direction changes. Harmful: toward-compliance drives non-refusal (jailbreak). Benign: pushing AWAY from compliance induces refusal of benign requests. Random (equal magnitude) does little in either until it degrades the model at large $|\alpha|$.