


Decoded == steered (goals[50:]): the compliance directions and the prior task-only refusal direction both steer compliance; the decodable harm direction won't move harm past ~1.6

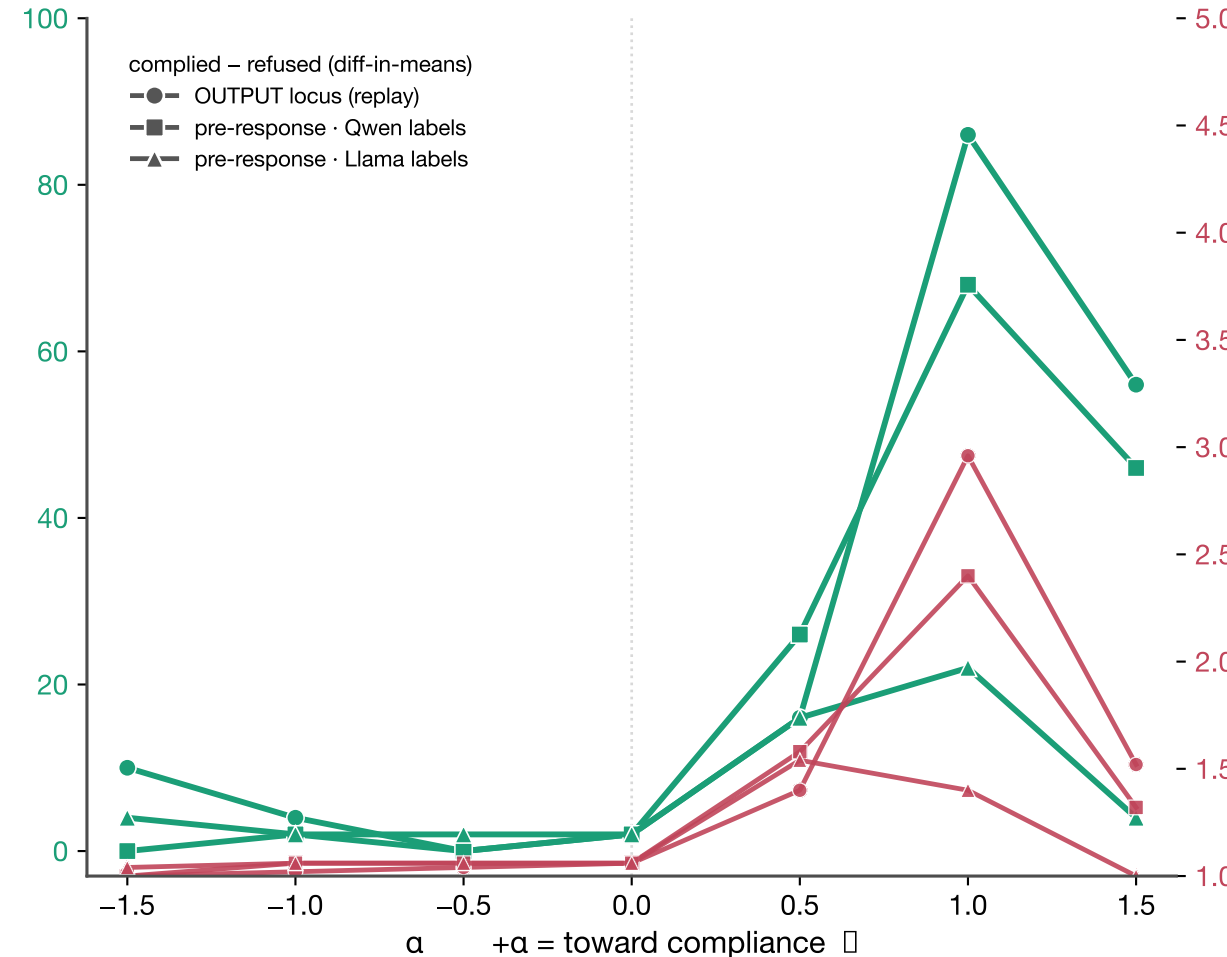
PRIOR, TASK-ONLY direction □ it STEERS

refusal_dm = $\mu(\text{harmful}) - \mu(\text{benign})$ PROMPTS · pre-response · a different contrast



COMPLIANCE directions □ they STEER

complied – refused (training corpus) · goals[50:] (n=50) · Llama judge



HARM directions □ none STEER harm

$\mu(\geq 4) - \mu(\leq 2)$ etc. · crimson=harm (bold), teal=comp% (faint) · goals[50:] · Llama

